

## Importing Libraries

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

## Importing Datasets

```
In [2]: df=pd.read_csv("innovation_and_development_database.csv")
df
```

Out[2]:

	country	code	year	eap	eca	lac	mena	sha	sa	hi	...	y	stockpatEPO	poptotal	labor	rdexpgdp	pat
0	Aruba	ABW	1960	0	0	1	0	0	0	0.0	...	NaN	NaN	NaN	NaN	NaN	
1	Aruba	ABW	1961	0	0	1	0	0	0	0.0	...	NaN	NaN	NaN	NaN	NaN	
2	Aruba	ABW	1962	0	0	1	0	0	0	0.0	...	NaN	NaN	NaN	NaN	NaN	
3	Aruba	ABW	1963	0	0	1	0	0	0	0.0	...	NaN	NaN	NaN	NaN	NaN	
4	Aruba	ABW	1964	0	0	1	0	0	0	0.0	...	NaN	NaN	NaN	NaN	NaN	
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
8290	Zimbabwe	ZWE	1998	0	0	0	0	1	0	0.0	...	8.290000e+09	8.0	12153850.0	6236700.0	NaN	
8291	Zimbabwe	ZWE	1999	0	0	0	0	1	0	0.0	...	8.230000e+09	8.0	12388320.0	6374300.0	NaN	
8292	Zimbabwe	ZWE	2000	0	0	0	0	1	0	0.0	...	7.830000e+09	8.0	12627000.0	6514800.0	NaN	
8293	Zimbabwe	ZWE	2001	0	0	0	0	1	0	0.0	...	NaN	8.0	12820650.0	NaN	NaN	
8294	Zimbabwe	ZWE	2002	0	0	0	0	1	0	0.0	...	NaN	NaN	NaN	NaN	NaN	

8295 rows × 33 columns



# Data Cleaning and Data Preprocessing

In [3]: `df=df.fillna(1)`  
`df`

Out[3]:

	country	code	year	eap	eca	lac	mena	sha	sa	hi	...	y	stockpatEPO	poptotal	labor	rdexpgdp	
0	Aruba	ABW	1960	0	0	1	0	0	0	0.0	...	1.000000e+00	1.0	1.0	1.0	1.0	
1	Aruba	ABW	1961	0	0	1	0	0	0	0.0	...	1.000000e+00	1.0	1.0	1.0	1.0	
2	Aruba	ABW	1962	0	0	1	0	0	0	0.0	...	1.000000e+00	1.0	1.0	1.0	1.0	
3	Aruba	ABW	1963	0	0	1	0	0	0	0.0	...	1.000000e+00	1.0	1.0	1.0	1.0	
4	Aruba	ABW	1964	0	0	1	0	0	0	0.0	...	1.000000e+00	1.0	1.0	1.0	1.0	
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	
8290	Zimbabwe	ZWE	1998	0	0	0	0	1	0	0.0	...	8.290000e+09	8.0	12153850.0	6236700.0	1.0	
8291	Zimbabwe	ZWE	1999	0	0	0	0	1	0	0.0	...	8.230000e+09	8.0	12388320.0	6374300.0	1.0	
8292	Zimbabwe	ZWE	2000	0	0	0	0	1	0	0.0	...	7.830000e+09	8.0	12627000.0	6514800.0	1.0	
8293	Zimbabwe	ZWE	2001	0	0	0	0	1	0	0.0	...	1.000000e+00	8.0	12820650.0	1.0	1.0	
8294	Zimbabwe	ZWE	2002	0	0	0	0	1	0	0.0	...	1.000000e+00	1.0	1.0	1.0	1.0	

In [4]: `df.columns`

Out[4]: Index(['country', 'code', 'year', 'eap', 'eca', 'lac', 'mena', 'sha', 'sa',  
'hi', 'pat', 'patepo', 'royal', 'rdexp', 'rdper', 'rdfinabro',  
'rdfinprod', 'rdperfprod', 'rdperfhe', 'rdperfpub', 'lowrdexp',  
'lowrdfinprod', 'lowrdperfprod', 'y', 'stockpatEPO', 'poptotal',  
'labor', 'rdexpgdp', 'patgrantedstock', 'plantpatstock',  
'designpatstock', 'plantpat', 'designpat'],  
dtype='object')

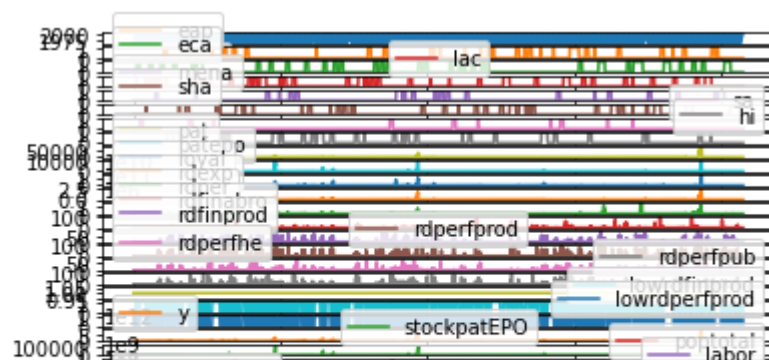
In [5]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8295 entries, 0 to 8294
Data columns (total 33 columns):
#   Column                Non-Null Count  Dtype
---  -
0   country               8295 non-null   object
1   code                  8295 non-null   object
2   year                  8295 non-null   int64
3   eap                   8295 non-null   int64
4   eca                   8295 non-null   int64
5   lac                   8295 non-null   int64
6   mena                  8295 non-null   int64
7   sha                   8295 non-null   int64
8   sa                    8295 non-null   int64
9   hi                    8295 non-null   float64
10  pat                   8295 non-null   float64
11  patepo                8295 non-null   float64
12  royal                 8295 non-null   float64
13  rdexp                 8295 non-null   float64
14  rdper                 8295 non-null   float64
15  rdfinabro             8295 non-null   float64
16  rdfinprod             8295 non-null   float64
17  rdperfprod            8295 non-null   float64
18  rdperfhe              8295 non-null   float64
19  rdperfpub             8295 non-null   float64
20  lowrdexp              8295 non-null   float64
21  lowrdfinprod          8295 non-null   float64
22  lowrdperfprod         8295 non-null   float64
23  y                     8295 non-null   float64
24  stockpatEPO           8295 non-null   float64
25  poptotal              8295 non-null   float64
26  labor                 8295 non-null   float64
27  rdexpgdp              8295 non-null   float64
28  patgrantedstock       8295 non-null   float64
29  plantpatstock         8295 non-null   float64
30  designpatstock        8295 non-null   float64
31  plantpat              8295 non-null   float64
32  designpat             8295 non-null   float64
dtypes: float64(24), int64(7), object(2)
memory usage: 2.1+ MB
```

# Line chart

In [6]: `df.plot.line(subplots=True)`


Out[6]: array([<AxesSubplot:>, <AxesSubplot:>, <AxesSubplot:>, <AxesSubplot:>, <AxesSubplot:>, <AxesSubplot:>, <AxesSubplot:>, <AxesSubplot:>, <AxesSubplot:>, <AxesSubplot:>, <AxesSubplot:>, <AxesSubplot:>, <AxesSubplot:>, <AxesSubplot:>, <AxesSubplot:>, <AxesSubplot:>, <AxesSubplot:>, <AxesSubplot:>, <AxesSubplot:>], dtype=object)



# Line chart

In [7]: `df.plot.line()`

Out[7]: <AxesSubplot:>




## Bar chart

In [8]: `b=df[0:50]`

```
In [9]: b.plot.bar()
```


Out[9]: <AxesSubplot:>



# Histogram

```
In [10]: df.plot.hist()
```


```
Out[10]: <AxesSubplot:ylabel='Frequency'>
```



## Area chart

```
In [11]: df.plot.area()
```

```
Out[11]: <AxesSubplot:>
```




## Box chart



```
In [12]: df.plot.box()
```

```
Out[12]: <AxesSubplot:>
```




## Pie chart

```
In [13]: b.plot.pie(y='designpat' )
```

```
Out[13]: <AxesSubplot:ylabel='designpat'>
```








## Scatter chart

```
In [14]: df.plot.scatter( x='rdexpgdp',y= 'patgrantedstock')
```

```
Out[14]: <AxesSubplot:xlabel='rdexpgdp', ylabel='patgrantedstock'>
```



In [15]: df.info()

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8295 entries, 0 to 8294
Data columns (total 33 columns):
#   Column                Non-Null Count  Dtype
---  -
0   country                8295 non-null   object
1   code                   8295 non-null   object
2   year                   8295 non-null   int64
3   eap                    8295 non-null   int64
4   eca                    8295 non-null   int64
5   lac                    8295 non-null   int64
6   mena                   8295 non-null   int64
7   sha                     8295 non-null   int64
8   sa                     8295 non-null   int64
9   hi                     8295 non-null   float64
10  pat                    8295 non-null   float64
11  patepo                 8295 non-null   float64
12  royal                  8295 non-null   float64
13  rdexp                  8295 non-null   float64
14  ...                     ...

```

In [16]: df.describe()

Out[16]:

	year	eap	eca	lac	mena	sha	sa	hi	pat	pi
<b>count</b>	8295.000000	8295.000000	8295.000000	8295.000000	8295.000000	8295.000000	8295.000000	8295.000000	8295.000000	8295.000000
<b>mean</b>	1981.203014	0.094515	0.195901	0.176974	0.098614	0.150090	0.026522	0.134901	411.663773	65.760000
<b>std</b>	12.421590	0.292561	0.396917	0.381670	0.298161	0.357182	0.160691	0.341638	3893.360772	580.340000
<b>min</b>	1960.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
<b>25%</b>	1970.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	1.000000
<b>50%</b>	1981.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	1.000000
<b>75%</b>	1992.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	2.000000	1.000000
<b>max</b>	2002.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	87607.000000	10300.000000

8 rows × 31 columns

```
In [17]: df1=df[['year', 'eap', 'eca', 'lac', 'mena', 'sha', 'sa',  
               'hi', 'pat', 'patepo', 'royal', 'rdexp', 'rdper', 'rdfinabro',  
               'rdfinprod', 'rdperfprod', 'rdperfhe', 'rdperfpub', 'lowrdexp',  
               'lowrdfinprod', 'lowrdperfprod', 'y', 'stockpatEPO', 'poptotal',  
               'labor', 'rdexpgdp', 'patgrantedstock', 'plantpatstock',  
               'designpatstock', 'plantpat', 'designpat']]
```


## EDA AND VISUALIZATION


```
In [18]: sns.pairplot(df1[0:50])
```

```
Out[18]: <seaborn.axisgrid.PairGrid at 0x24b293635e0>
```









```
In [19]: sns.distplot(df1['designpat'])
```


C:\Users\USER\anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).  
warnings.warn(msg, FutureWarning)

```
Out[19]: <AxesSubplot:xlabel='designpat', ylabel='Density'>
```



```
In [20]: sns.heatmap(df1.corr())
```

```
Out[20]: <AxesSubplot:>
```



## TO TRAIN THE MODEL AND MODEL BUILDING

```
In [21]: x=df[['year', 'eap', 'eca', 'lac', 'mena', 'sha', 'sa',
               'hi', 'pat', 'patepo', 'royal', 'rdexp', 'rdper', 'rdinabro',
               'rdinprod', 'rdperfprod', 'rdperfh', 'rdperfpub', 'lowrdexp',
               'lowrdinprod', 'lowrdperfprod', 'y', 'stockpatEPO', 'poptotal',
               'labor', 'rdexpdp', 'patgrantedstock', 'plantpatstock',
               'designpatstock', 'plantpat']]
y=df['designpat']
```

```
In [22]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

# Linear Regression

```
In [23]: from sklearn.linear_model import LinearRegression  
lr=LinearRegression()  
lr.fit(x_train,y_train)
```

Out[23]: LinearRegression()

```
In [24]: lr.intercept_
```

Out[24]: -895.707237255563

```
In [25]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])  
coeff
```


Out[25]:

	Co-efficient
<b>year</b>	4.478269e-01
<b>eap</b>	1.044973e+01
<b>eca</b>	-4.112414e-01
<b>lac</b>	1.711691e+00
<b>mena</b>	1.505925e+00
<b>sha</b>	2.727725e+00
<b>sa</b>	2.171481e+00
<b>hi</b>	-1.068480e+01
<b>pat</b>	1.077552e-02
<b>patepo</b>	-2.062311e-02
<b>royal</b>	-1.515536e-08
<b>rdexp</b>	-2.591489e-09
<b>rdper</b>	-5.410640e-05
<b>rdfinabro</b>	-1.926877e-01
<b>rdfinprod</b>	3.072198e-01
<b>rdperfprod</b>	2.148397e-01
<b>rdperfhe</b>	-1.150555e-01
<b>rdperfpub</b>	5.812949e-03
<b>lowrdexp</b>	5.284662e-14
<b>lowrdfinprod</b>	2.735943e+00
<b>lowrdperfprod</b>	7.338671e-01
<b>y</b>	3.814805e-11
<b>stockpatEPO</b>	9.898542e-04
<b>poptotal</b>	-3.071126e-08
<b>labor</b>	3.423496e-08
<b>rdexpgdp</b>	6.661779e-01

Co-efficient	
<b>patgrantedstock</b>	-6.384360e-05
<b>plantpatstock</b>	-1.126773e+00
<b>designpatstock</b>	9.806644e-02
<b>plantpat</b>	1.316892e+01

```
In [26]: prediction =lr.predict(x_test)  
plt.scatter(y_test,prediction)
```

```
Out[26]: <matplotlib.collections.PathCollection at 0x24b5f465730>
```



## ACCURACY

```
In [27]: lr.score(x_test,y_test)
```

```
Out[27]: 0.9359440854000607
```

```
In [28]: lr.score(x_train,y_train)
```

```
Out[28]: 0.9556190976958261
```



## Ridge and Lasso

```
In [29]: from sklearn.linear_model import Ridge,Lasso
```

```
In [30]: rr=Ridge(alpha=10)  
rr.fit(x_train,y_train)
```

C:\Users\USER\anaconda3\lib\site-packages\sklearn\linear\_model\\_ridge.py:157: LinAlgWarning: Ill-conditioned matrix (rcond=6.91987e-27): result may not be accurate.  
return linalg.solve(A, Xy, sym\_pos=True, overwrite\_a=True).T

```
Out[30]: Ridge(alpha=10)
```

## Accuracy(Ridge)

```
In [31]: rr.score(x_test,y_test)
```

```
Out[31]: 0.9359433503141961
```

```
In [32]: rr.score(x_train,y_train)
```

```
Out[32]: 0.9556189968019976
```

```
In [33]: la=Lasso(alpha=10)  
la.fit(x_train,y_train)
```

```
Out[33]: Lasso(alpha=10)
```

```
In [34]: la.score(x_train,y_train)
```

```
Out[34]: 0.9553951543156738
```

## Accuracy(Lasso)

```
In [35]: la.score(x_test,y_test)
```

```
Out[35]: 0.9365057709302195
```

## ElasticNet

```
In [36]: from sklearn.linear_model import ElasticNet  
en=ElasticNet()  
en.fit(x_train,y_train)
```

```
Out[36]: ElasticNet()
```

```
In [37]: en.coef_
```

```
Out[37]: array([ 4.80389038e-01,  5.81402950e-01, -0.00000000e+00,  0.00000000e+00,  
                -0.00000000e+00,  0.00000000e+00, -0.00000000e+00, -6.77890367e-01,  
                 1.09871047e-02, -1.95517371e-02, -1.58117666e-08, -2.30292806e-09,  
                -5.39439272e-05, -1.70086963e-01,  2.00042804e-01,  2.29682240e-01,  
                -1.76965155e-01,  0.00000000e+00,  0.00000000e+00,  0.00000000e+00,  
                 0.00000000e+00,  2.79875266e-11,  8.57046944e-04, -2.65239954e-08,  
                 4.22220894e-08, -0.00000000e+00, -5.45100706e-05, -1.11854094e+00,  
                 9.80471935e-02,  1.30136367e+01])
```

```
In [38]: en.intercept_
```

```
Out[38]: -957.1590613062978
```

```
In [39]: prediction=en.predict(x_test)
```

```
In [40]: en.score(x_test,y_test)
```

```
Out[40]: 0.9361560503753533
```

## Evaluation Metrics

```
In [41]: from sklearn import metrics
print(metrics.mean_absolute_error(y_test,prediction))
print(metrics.mean_squared_error(y_test,prediction))
print(np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

```
13.24769486414481
10640.218004441078
103.15143239161091
```

## Logistic Regression

```
In [42]: from sklearn.linear_model import LogisticRegression
```

```
In [43]: feature_matrix=df[['year', 'eap', 'eca', 'lac', 'mena', 'sha', 'sa',
    'hi', 'pat', 'patepo', 'royal', 'rdexp', 'rdper', 'rdfinabro',
    'rdfinprod', 'rdperfprod', 'rdperfhe', 'rdperfpub', 'lowrdexp',
    'lowrdfinprod', 'lowrdperfprod', 'y', 'stockpatEPO', 'poptotal',
    'labor', 'rdexpgdp', 'patgrantedstock', 'plantpatstock',
    'designpatstock', 'plantpat']]
target_vector=df[ 'designpat']
```

```
In [44]: feature_matrix.shape
```

```
Out[44]: (8295, 30)
```

```
In [45]: target_vector.shape
```

```
Out[45]: (8295,)
```

```
In [46]: from sklearn.preprocessing import StandardScaler
```

```
In [47]: fs=StandardScaler().fit_transform(feature_matrix)
```

```
In [48]: logr=LogisticRegression(max_iter=10000)
logr.fit(fs,target_vector)
```

```
Out[48]: LogisticRegression(max_iter=10000)
```

```
In [49]: observation=[[1,2,3,4,5,6,7,8,9,10,11,12,13,14,15,16,17,18,19,20,21,22,23,24,25,26,27,28,29,30]]
```

```
In [50]: prediction=logr.predict(observation)
print(prediction)
```

```
[212.]
```

In [51]: `logr.classes_`

```
Out[51]: array([0.0000e+00, 1.0000e+00, 2.0000e+00, 3.0000e+00, 4.0000e+00,
5.0000e+00, 6.0000e+00, 7.0000e+00, 8.0000e+00, 9.0000e+00,
1.0000e+01, 1.1000e+01, 1.2000e+01, 1.3000e+01, 1.4000e+01,
1.5000e+01, 1.6000e+01, 1.7000e+01, 1.8000e+01, 1.9000e+01,
2.0000e+01, 2.1000e+01, 2.2000e+01, 2.3000e+01, 2.4000e+01,
2.5000e+01, 2.6000e+01, 2.7000e+01, 2.8000e+01, 2.9000e+01,
3.0000e+01, 3.2000e+01, 3.3000e+01, 3.4000e+01, 3.7000e+01,
3.8000e+01, 3.9000e+01, 4.0000e+01, 4.1000e+01, 4.2000e+01,
4.3000e+01, 4.4000e+01, 4.5000e+01, 4.6000e+01, 4.7000e+01,
4.8000e+01, 4.9000e+01, 5.0000e+01, 5.4000e+01, 5.5000e+01,
5.6000e+01, 5.7000e+01, 5.8000e+01, 5.9000e+01, 6.0000e+01,
6.1000e+01, 6.2000e+01, 6.3000e+01, 6.4000e+01, 6.5000e+01,
6.6000e+01, 6.7000e+01, 6.9000e+01, 7.0000e+01, 7.1000e+01,
7.2000e+01, 7.3000e+01, 7.4000e+01, 7.6000e+01, 7.7000e+01,
7.8000e+01, 8.0000e+01, 8.1000e+01, 8.2000e+01, 8.3000e+01,
8.4000e+01, 8.5000e+01, 8.6000e+01, 8.7000e+01, 8.8000e+01,
8.9000e+01, 9.0000e+01, 9.1000e+01, 9.3000e+01, 9.4000e+01,
9.5000e+01, 9.6000e+01, 9.7000e+01, 9.8000e+01, 9.9000e+01,
1.0000e+02, 1.0100e+02, 1.0200e+02, 1.0300e+02, 1.0500e+02,
1.0600e+02, 1.0700e+02, 1.0900e+02, 1.1000e+02, 1.1200e+02,
1.1300e+02, 1.1400e+02, 1.1500e+02, 1.1700e+02, 1.1800e+02,
1.1900e+02, 1.2000e+02, 1.2100e+02, 1.2300e+02, 1.2400e+02,
1.2500e+02, 1.2600e+02, 1.2700e+02, 1.2900e+02, 1.3200e+02,
1.3300e+02, 1.3400e+02, 1.3600e+02, 1.3800e+02, 1.3900e+02,
1.4200e+02, 1.4300e+02, 1.4400e+02, 1.4700e+02, 1.5000e+02,
1.5200e+02, 1.5300e+02, 1.5600e+02, 1.5900e+02, 1.6000e+02,
1.6200e+02, 1.6300e+02, 1.6500e+02, 1.6900e+02, 1.7200e+02,
1.7300e+02, 1.7600e+02, 1.7900e+02, 1.8000e+02, 1.8100e+02,
1.8200e+02, 1.8400e+02, 1.8500e+02, 1.8600e+02, 1.9000e+02,
1.9300e+02, 1.9500e+02, 1.9600e+02, 1.9700e+02, 2.0100e+02,
2.0200e+02, 2.0500e+02, 2.0800e+02, 2.1100e+02, 2.1200e+02,
2.1300e+02, 2.1500e+02, 2.1600e+02, 2.2100e+02, 2.2200e+02,
2.2700e+02, 2.2800e+02, 2.3000e+02, 2.3100e+02, 2.3400e+02,
2.3700e+02, 2.3900e+02, 2.4000e+02, 2.4300e+02, 2.4700e+02,
2.5000e+02, 2.5300e+02, 2.5400e+02, 2.5700e+02, 2.5800e+02,
2.6000e+02, 2.6500e+02, 2.7500e+02, 2.8200e+02, 3.0000e+02,
3.0600e+02, 3.2000e+02, 3.3000e+02, 3.3800e+02, 3.4100e+02,
3.5000e+02, 3.5600e+02, 3.6000e+02, 3.6800e+02, 3.7000e+02,
3.7200e+02, 3.8200e+02, 3.9000e+02, 3.9600e+02, 4.0100e+02,
4.1000e+02, 4.1800e+02, 4.3800e+02, 4.3900e+02, 4.6600e+02,
4.8200e+02, 4.8400e+02, 4.8500e+02, 5.0300e+02, 5.0500e+02,
5.0900e+02, 5.2200e+02, 5.3900e+02, 5.4700e+02, 5.7600e+02,
5.8800e+02, 6.2000e+02, 6.9800e+02, 7.0500e+02, 7.9500e+02,
```

```
8.3300e+02, 8.6100e+02, 9.3600e+02, 9.4800e+02, 1.0260e+03,  
1.0370e+03, 1.0560e+03, 1.1350e+03, 1.1370e+03, 1.1490e+03,  
1.1680e+03, 1.2070e+03, 1.2940e+03, 1.3070e+03, 1.3100e+03,  
1.3640e+03, 1.4970e+03, 1.5460e+03, 2.4690e+03, 3.0520e+03,  
3.0550e+03, 3.0650e+03, 3.2780e+03, 3.4280e+03, 3.4460e+03,  
3.4750e+03, 3.5460e+03, 3.5700e+03, 3.6450e+03, 3.8830e+03,  
3.9020e+03, 5.0690e+03, 6.0130e+03, 6.0750e+03, 7.4160e+03,  
7.6970e+03, 7.7470e+03, 7.8630e+03, 8.2510e+03, 9.3250e+03,  
9.6540e+03, 9.9130e+03, 1.0346e+04, 1.1285e+04])
```

```
In [52]: logr.score(fs,target_vector)
```

```
Out[52]: 0.8764315852923448
```

```
In [53]: logr.predict_proba(observation)[0][0]
```

```
Out[53]: 0.0
```

```
In [54]: logr.predict_proba(observation)
```



```
Out[54]: array([[0.00000000e+00, 6.46875839e-85, 4.18064698e-86, 4.97640492e-80,  
8.59945817e-83, 2.37734566e-56, 3.98520696e-68, 2.90827307e-74,  
3.93470301e-68, 7.84564082e-60, 4.43221229e-85, 6.56263221e-55,  
1.03039805e-73, 1.08013371e-59, 1.08112169e-50, 7.94176532e-68,  
1.81569452e-50, 2.60212594e-61, 4.31128405e-50, 7.36175743e-49,  
3.51381920e-52, 8.03683141e-50, 4.61455481e-44, 9.29966791e-71,  
2.15616849e-32, 1.47123590e-51, 3.22397464e-25, 1.13656197e-42,  
3.84756439e-36, 3.92528441e-58, 2.57777087e-38, 4.08345965e-33,  
4.43714065e-49, 1.95954619e-43, 1.45852110e-54, 1.32840786e-50,  
1.61612437e-22, 2.16834024e-29, 1.26908838e-34, 8.98132380e-30,  
2.17424290e-44, 4.15451744e-15, 6.05958093e-36, 3.76411606e-12,  
1.64039866e-26, 5.78586665e-31, 1.98698848e-18, 1.60937793e-35,  
3.86587379e-25, 7.65725299e-09, 1.68641093e-36, 7.44871984e-19,  
9.91118520e-31, 3.27302787e-26, 4.54761528e-32, 3.46546939e-26,  
3.43068716e-38, 2.09551086e-46, 1.47074020e-24, 5.60652672e-29,  
4.46188339e-14, 5.23448116e-35, 4.90515430e-13, 4.19574969e-43,  
1.88139168e-38, 5.87960896e-41, 1.68589460e-22, 1.06641334e-28,  
4.31086801e-32, 2.11026858e-21, 1.36350787e-13, 2.68466590e-17,  
2.06256309e-27, 4.64149327e-37, 6.11020490e-11, 1.77915082e-04,  
1.04327068e-33, 9.27239443e-25, 1.64059124e-19, 1.41867503e-26,  
3.45576309e-25, 3.78644588e-25, 5.53565304e-21, 1.10810312e-32,  
4.17326226e-30, 2.85198531e-14, 1.70870740e-16, 2.37696717e-31,  
4.57363954e-20, 1.10518235e-31, 7.04942289e-18, 7.09899878e-15,  
6.18243195e-20, 5.14010433e-12, 1.30131822e-26, 2.75215467e-28,  
3.56175114e-23, 1.43488680e-20, 1.49026166e-35, 2.31955497e-19,  
2.80943383e-28, 3.26844661e-25, 1.98767004e-25, 2.90381563e-47,  
4.16144030e-38, 2.15517707e-26, 5.16987221e-29, 3.01843396e-43,  
2.43209628e-24, 1.67542453e-22, 1.49536949e-44, 4.37752030e-24,  
1.19030241e-27, 1.26335802e-12, 3.25590978e-18, 1.46940286e-36,  
6.21194731e-25, 1.02987622e-16, 2.55264374e-31, 3.08002770e-28,  
2.73592092e-19, 1.09563396e-20, 7.28661834e-15, 8.63266119e-24,  
6.55186889e-19, 3.67493263e-23, 5.37165816e-11, 1.42090765e-29,  
1.96504387e-26, 3.15806890e-35, 4.67423646e-41, 1.03249746e-42,  
2.06352084e-18, 4.80872013e-11, 2.38144854e-22, 1.44623746e-19,  
1.12561121e-23, 1.27398771e-17, 5.88348615e-41, 9.14398660e-13,  
2.69631717e-12, 3.53809810e-06, 1.35670563e-37, 4.00635816e-07,  
3.66333929e-26, 3.21990721e-11, 4.34472713e-18, 3.82940747e-11,  
4.48980290e-12, 3.70126759e-10, 3.42435599e-16, 2.23506481e-29,  
2.06547055e-22, 1.65143853e-23, 9.99292856e-01, 1.89931562e-40,  
9.95178307e-09, 4.04014787e-33, 6.45046777e-19, 6.39481242e-17,  
1.65711204e-21, 1.66902668e-13, 5.85102749e-09, 3.58553149e-20,  
7.63863608e-14, 2.16528932e-16, 2.31182283e-19, 5.19372698e-10,  
2.85136582e-18, 2.20030544e-23, 8.40363376e-15, 1.52062520e-10,
```

```
2.11702071e-34, 9.43588493e-11, 2.38638165e-34, 6.84723725e-16,  
7.27402502e-24, 1.34617982e-10, 7.45356898e-05, 2.01297770e-13,  
6.90786546e-16, 2.47182409e-26, 4.93804370e-15, 1.42551905e-13,  
2.26831073e-18, 4.95819552e-17, 2.54256226e-26, 3.33888563e-11,  
1.20758731e-31, 2.15710411e-24, 4.34686023e-07, 1.48682527e-40,  
1.80994770e-16, 1.56738491e-25, 5.46790333e-18, 3.54425763e-04,  
1.27586197e-10, 3.05700233e-12, 1.93730817e-21, 5.51384567e-23,  
2.69380119e-17, 3.77031259e-10, 1.01278030e-07, 5.00823656e-20,  
9.23034863e-05, 1.50877311e-19, 3.49094704e-21, 2.04317551e-20,  
6.66023049e-12, 2.20980382e-22, 5.86801028e-22, 8.35022702e-20,  
5.26476071e-19, 2.53780151e-18, 6.88038395e-17, 1.31007494e-15,  
3.18385328e-17, 1.42025828e-18, 3.67892414e-17, 2.01043312e-14,  
6.10377056e-18, 5.68120253e-21, 2.01399099e-13, 7.29833544e-15,  
3.86518040e-15, 6.63952897e-26, 9.64773366e-12, 1.62949403e-24,  
9.21556898e-19, 2.45420341e-12, 5.20103815e-26, 1.31984657e-10,  
6.15952510e-20, 9.20354758e-19, 5.84975126e-21, 1.31131770e-18,  
8.45741933e-13, 1.03629389e-08, 1.79179668e-13, 2.82884116e-08,  
1.24337025e-11, 6.93118767e-12, 2.24629133e-13, 5.49609681e-07,  
6.10191031e-09, 9.96528220e-11, 3.87985456e-13, 2.53836642e-09,  
1.26770832e-06, 1.24209791e-08, 3.02700348e-08, 1.53830151e-06,  
2.39069364e-11, 2.79232187e-11, 2.18150800e-21, 2.44157090e-13,  
1.53833366e-08, 5.51628248e-12, 1.97159115e-09]])
```

## Random Forest

```
In [55]: from sklearn.ensemble import RandomForestClassifier
```

```
In [56]: rfc=RandomForestClassifier()  
rfc.fit(x_train,y_train)
```

```
Out[56]: RandomForestClassifier()
```

```
In [57]: parameters={'max_depth':[1,2,3,4,5],  
                    'min_samples_leaf':[5,10,15,20,25],  
                    'n_estimators':[10,20,30,40,50]  
}
```

```
In [58]: from sklearn.model_selection import GridSearchCV
grid_search = GridSearchCV(estimator=rfc, param_grid=parameters, cv=2, scoring="accuracy")
grid_search.fit(x_train, y_train)
```

C:\Users\USER\anaconda3\lib\site-packages\sklearn\model\_selection\\_split.py:676: UserWarning: The least populated class in y has only 1 members, which is less than n\_splits=2.  
warnings.warn(

```
Out[58]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
                    param_grid={'max_depth': [1, 2, 3, 4, 5],
                                'min_samples_leaf': [5, 10, 15, 20, 25],
                                'n_estimators': [10, 20, 30, 40, 50]},
                    scoring='accuracy')
```

```
In [59]: grid_search.best_score_
```

```
Out[59]: 0.8815018945918016
```

```
In [60]: rfc_best = grid_search.best_estimator_
```



```
In [62]: print("Linear Regression:",lr.score(x_test,y_test))  
print("Ridge Regression:",rr.score(x_test,y_test))  
print("Lasso Regression",la.score(x_test,y_test))  
print("ElasticNet Regression:",en.score(x_test,y_test))  
print("Logistic Regression:",logr.score(fs,target_vector))  
print("Random Forest:",grid_search.best_score_)
```

Linear Regression: 0.9359440854000607  
Ridge Regression: 0.9359433503141961  
Lasso Regression 0.9365057709302195  
ElasticNet Regression: 0.9361560503753533  
Logistic Regression: 0.8764315852923448  
Random Forest: 0.8815018945918016

## Lasso Regression is suitable for this dataset